

Case study: test-set trajectories with given (k-means) vs predicted (tabmixer) labels
test accuracy 99.99% — red title/border = misclassified

